

# CSCI B 565 Data Mining Bonus Homework

Ganesh Nagarajan, [gnagaraj@indiana.edu](mailto:gnagaraj@indiana.edu)

December 17, 2015

All work herein is solely mine.

## Solutions

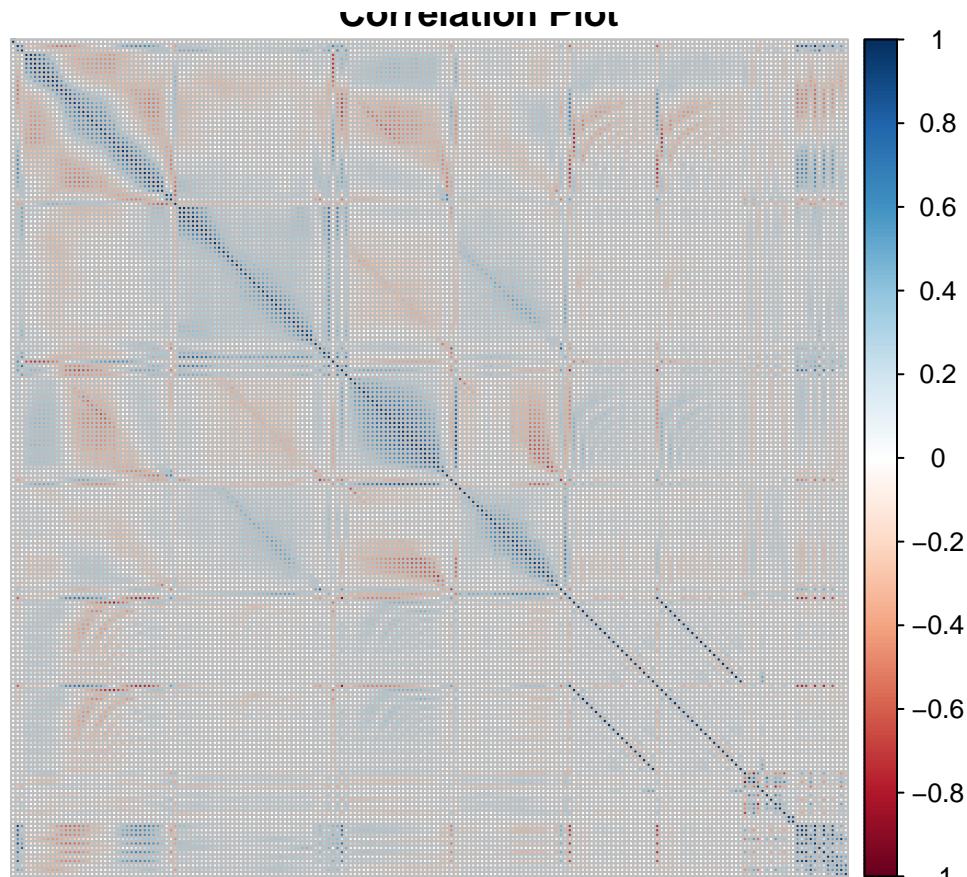
1. Relative Frequencies of these labels

```
# Load the Train Dataset to R
tuned <- read.csv("genresTrain.csv")
# Create a Frequency Table with the Genre
table(tuned[,192])
```

```
##
##      Blues Classical      Jazz      Metal      Pop      Rock
##      1596         3444      3003       924      1575      1953
```

2. Correlation of these features

```
library(corrplot)
num.tuned <- tuned[,192]
# Find Correlation
cor.tune <- cor(num.tuned)
# Plot Correlation
corrplot(cor.tune, method = c("square"), title = "Correlation Plot", tl.pos = c("n"))
```



```
cor.upper.tune <- cor.tune
# Make the Lower triangular matrix of the correlation matrix 0
cor.upper.tune[lower.tri(cor.upper.tune)] <- 0
# Since the diagonals are always correlated, make them 0
diag(cor.upper.tune) <- 0
#Find the labels which have a perfect correlation
cno <- which(cor.upper.tune == 1,arr.ind = TRUE)
print("Elements which have perfect correlation")
```

```
## [1] "Elements which have perfect correlation"
```

```
cno
```

```
##          row col
## PAR_MFCC2 129 149
## PAR_MFCC3 130 150
## PAR_MFCC4 131 151
## PAR_MFCC5 132 152
## PAR_MFCC6 133 153
## PAR_MFCC7 134 154
## PAR_MFCC8 135 155
## PAR_MFCC9 136 156
## PAR_MFCC10 137 157
## PAR_MFCC11 138 158
## PAR_MFCC12 139 159
```

```
## PAR_MFCC15 142 162
## PAR_MFCC16 143 163
## PAR_MFCC17 144 164
## PAR_MFCC18 145 165
## PAR_MFCC19 146 166
## PAR_MFCC20 147 167
```

```
cno1 <- which(cor.upper.tune > 0.9,arr.ind = TRUE)
print("Elements which have correlation with > 0.9")
```

```
## [1] "Elements which have correlation with > 0.9"
```

```
cno1
```

```
##          row col
## PAR_ASE1      4  5
## PAR_ASE2      5  6
## PAR_ASE3      6  7
## PAR_ASE4      7  8
## PAR_ASE5      8  9
## PAR_ASE6      9 10
## PAR_ASE29    32 33
## PAR_ASE30    33 34
## PAR_ASEV1    39 40
## PAR_ASEV2    40 41
## PAR_ASEV3    41 42
## PAR_ASEV4    42 43
## PAR_ASEV5    43 44
## PAR_SFM11    88 89
## PAR_SFM12    89 90
## PAR_SFM13    90 91
## PAR_SFM14    91 92
## PAR_SFM15    92 93
## PAR_SFM18    95 96
## PAR_SFM19    96 97
## PAR_MFCC1   128 148
## PAR_MFCC2   129 149
## PAR_MFCC3   130 150
## PAR_MFCC4   131 151
## PAR_MFCC5   132 152
## PAR_MFCC6   133 153
## PAR_MFCC7   134 154
## PAR_MFCC8   135 155
## PAR_MFCC9   136 156
## PAR_MFCC10  137 157
## PAR_MFCC11  138 158
## PAR_MFCC12  139 159
## PAR_MFCC13  140 160
## PAR_MFCC14  141 161
## PAR_MFCC15  142 162
## PAR_MFCC16  143 163
## PAR_MFCC17  144 164
## PAR_MFCC18  145 165
```

```
## PAR_MFCC19    146 166
## PAR_MFCC20    147 167
## PAR_SC        2 180
## PAR_SC        2 184
## PAR_ZCD       180 184
## PAR_1RMS_TCD 181 186
## PAR_2RMS_TCD 182 188
## PAR_3RMS_TCD 183 190
```

- It can be seen that about 47 of these features are correlated. An analysis is done in the next section considering and omitting these features.

3. There were about 60 performers
4. 15-20 pieces were performed by each performer
5. The first segment of the 20 segments with 191 features are available.
6. Building the Classifier

- **Building a Naive Bayes Classifier with Feature Engineering and Normalization of the Data**
- A Naive bayes classifier was implemented for the training set. As discussed in the previous section, there are about 47 correlated features. This section removes these correlated features except one and the data is normalized by using the standardization formula,  $\hat{x} = \frac{x - \bar{x}}{\sigma}$

```
# Remove the correlated columns
tune.ds <- tuned[c(cno[,2],cno1[,2])*-1]

#Normalization function
normalize <- function(df){
  v_cols<-ncol(df)
  df_normalized <- df
  for (i in 1:v_cols){
    df_normalized[,i]<-(df_normalized[,i]-mean(df_normalized[,i]))/sd(df_normalized[,i])
  }
  return(df_normalized)
}

library(e1071)

#Normalize the dataset
df_norm.tune <- normalize(tune.ds[-147])
df_norm.tune["GENRE"]<-tune.ds[147]

# Build a Naive Bayes Model
s<-naiveBayes(GENRE~.,data=df_norm.tune)
# Create the confusion Matrix for Train dataset
table(predict(s,df_norm.tune),df_norm.tune[,147])
```

```
##
##          Blues Classical Jazz Metal  Pop Rock
## Blues      1296         21  382    51   82  158
## Classical   43        3131 1070   12   22   34
## Jazz        13         267 1226    9   38   39
## Metal       104          0   26  766   22  349
```

```
##      Pop          61          13  140      39 1314  109
##      Rock          79          12  159      47   97 1264
```

```
# Load the Test Dataset
trainds <- read.csv("genresTest.csv")
# Remove same correlated columns
c_train.ds <- trainds[,c(cno[,2],cno1[,2])*-1]
# Normalize the test dataset
df_normal<-normalize(c_train.ds)
# Read Base line from KNN
baseline<-read.csv("genresBaseline.txt")
# Confusion Matrix between the Naive Bayes prediction and the baseline KNN prediction.
table(predict(s,df_normal[-1,]),baseline[,1])
```

```
##
##          Blues Classical Jazz Metal  Pop Rock
## Blues          3          0   2    0    2   0
## Classical      29         167  93    8   35  23
## Jazz           88        2739 1664   4   28  51
## Metal          188          3  42   109  11 118
## Pop            266          75 562   40 707 303
## Rock           269         146 509   40  95 1849
```

- Building a Naive Bayes Classifier with Feature Engineering
- This test run is without normalizing the data,

```
library(corrplot)
tunedds <- read.csv("genresTrain.csv")
num.tunedds<- tunedds[-192]
cor.tune <- cor(num.tunedds)
cor.upper.tune <- cor.tune
cor.upper.tune[lower.tri(cor.upper.tune)] <- 0
diag(cor.upper.tune) <- 0
cno <- which(cor.upper.tune == 1,arr.ind = TRUE)
cno1 <- which(cor.upper.tune > 0.9,arr.ind = TRUE)
# Removing correlations
tune.ds <- tunedds[,c(cno[,2],cno1[,2])*-1]

library(e1071)
# Train the Model
s<-naiveBayes(GENRE~.,data=tune.ds)
# Confusion Matrix of the Training set
table(predict(s,tune.ds),tune.ds[,147])
```

```
##
##          Blues Classical Jazz Metal  Pop Rock
## Blues       1311          21 382   51   88 158
## Classical    25        3131 1070   12  13  34
## Jazz         12        267 1227    9  28  39
## Metal        104          0  25   766  22 349
## Pop          65          13 140   39 1327 109
## Rock         79          12 159   47   97 1264
```

```

trainds <- read.csv("genresTest.csv")
c_train.ds <- trainds[,c(cno[,2],cno1[,2])*-1]
baseline<-read.csv("genresBaseline.txt")
# COnfusion matrix for the Naive Bayes classifier output and the KNN output.
table(predict(s,c_train.ds[-1,]),baseline[,1])

```

```

##
##          Blues Classical Jazz Metal  Pop Rock
## Blues      616         62  360    75  203  202
## Classical   11        2418  540     3    6   25
## Jazz        13         618 1635     4   24   54
## Metal       60          0   13    74    0   66
## Pop         85         20  169    16  569  254
## Rock        58         12  155    29   76 1743

```

- Building a Naive Bayes Classifier straight from data
- No feature engineering or standardization

```

library(corrplot)
tunedds <- read.csv("genresTrain.csv")

library(e1071)
s<-naiveBayes(GENRE~.,data=tunedds)
# Confusion Matrix for Train Dataset
table(predict(s,tunedds),tunedds[,192])

```

```

##
##          Blues Classical Jazz Metal  Pop Rock
## Blues     1311         14  403    39   74  160
## Classical   31        3133  977    10   23   16
## Jazz        14         278 1261    11   38   51
## Metal       86          0   19   776   24  339
## Pop         58          8  165    37 1319   94
## Rock        96         11  178    51   97 1293

```

```

trainds <- read.csv("genresTest.csv")
baseline<-read.csv("genresBaseline.txt")
# COnfusion Matrix for the Naive Bayes Duput and Knn baseline output
table(predict(s,trainds[-1,]),baseline[,1])

```

```

##
##          Blues Classical Jazz Metal  Pop Rock
## Blues      581         45  403    72  189  209
## Classical    8        2362  416     4   15   20
## Jazz        17         702 1716     8   32   56
## Metal       55          0   13    73    0   73
## Pop        108         14  162    12  577  198
## Rock        74          7  162    32   65 1788

```

## 7. Conclusions

- Standardization or normalization, atleast with this example created lot of noise and aberrant results.
- Removing correlated columns had almost same accuracy as the full dataset. Infact, the correlated columns removed dataset had more matching to blues for the output of the KNN than the full dataset.
- Matching the full dataset has the most match with the KNN baseline model.

Regarding the quality of the model, the full blown naive bayes has 7097 values classified out of 10269 values of test dataset. This brings up the accutacy of 69%.

8. An 1:NN algorithm was used for base solution. The proposed model is compared with the results of the baseline soution in the above question. ###References and acknowledgements
9. Packages Corrplot 0.73 and e1071 1.6-7 packages were used for this assignment.